

Machine Learning

State Space Search

---

---

---

---

---

---

---

State Space

Matt Johnson, Ph.D.2

---

---

---

---

---

---

---

State

A **state** of a search problem is an abstract representation of all relevant problem *features*.

A state can be viewed one possible step to take in the search.

- May not be necessary to examine each state
- May not be possible to get to each state

Matt Johnson, Ph.D.3

---

---

---

---

---

---

---

## State Space

The **state space** of a problem is the set of all possible states of the problem.

A state space can be viewed as a directed graph where each state is a node and each edge is an action that can be performed.

A state space can be finite or infinite.

Matt Johnson, Ph.D.

4

---

---

---

---

---

---

---

## State Space Search

**State space search** is a process in which successive configurations or states of an instance are considered, with the intention of finding a goal state with a desired property.

Matt Johnson, Ph.D.

5

---

---

---

---

---

---

---

## Search Problems

A **search problem** can be defined formally by 5 components:

1. States
2. Initial State
3. Successor Function
4. Goal Test
5. Path Cost

Matt Johnson, Ph.D.

6

---

---

---

---

---

---

---

## Initial State

The **initial state** of a problem is where the search agent begins.

- Specially designated
- Can be viewed as the root of a search tree

Matt Johnson, Ph.D.

7

---

---

---

---

---

---

---

## Successor Function

An *action* is how the agent moves between states.

A **successor function** applies each possible action to the current state

- Generates a set of reachable states
- May create states which have already been examined

Matt Johnson, Ph.D.

8

---

---

---

---

---

---

---

## Path Cost

The **path cost** is the calculated weight of the edges connecting two nodes in the state space.

- For many simple problems this is just the number of nodes in the path
- Cost may be an estimation

The **step cost** of an action is the path cost between two adjacent nodes when performing an action.

Matt Johnson, Ph.D.

9

---

---

---

---

---

---

---

## Goal Test

The **goal test** determines whether a given state is the goal state.

A **goal state** is a state that satisfies the goal test.

The goal test determines whether the current state:

- meets some abstract property
- is a member of an explicit state set

Matt Johnson, Ph.D.

10

---

---

---

---

---

---

---

## 8-Puzzle

### State:

the location of tiles numbered 1 through 8 and the blank tile in a 3\*3 grid

### Initial State:

an arbitrary state in the state space

1	4	3
7		6
5	8	2

Matt Johnson, Ph.D.

11

---

---

---

---

---

---

---

## 8-Puzzle (2)

### Successor Function:

legal states formed by moving tiles adjacent to the blank *up, down, left* and/or *right*

### Goal Test:

1	2	3
4	5	6
7	8	

### Path Cost:

Number of moves taken

Matt Johnson, Ph.D.

12

---

---

---


---

---

---

---

## 8-Puzzle (3)



Matt Johnson, Ph.D.

13

---

---

---

---

---

---

---

---

## 8-Queens

### States:

Any arrangement of 0-8 queens on a chess board

### Initial State:

An empty chess board

### Successor Function:

All possible chess boards obtained by adding a queen to any empty space

Matt Johnson, Ph.D.

14

---

---

---

---

---

---

---

---

## 8-Queens (2)

### Goal Test:

A chess board with 8 queens with no queen able to "attack" another queen on the board

### Path Cost:

Number of queens placed (irrelevant)

Matt Johnson, Ph.D.

15

---

---

---

---

---

---

---

---

# Routing Problem

**States:**

A location on the map, and the current time

**Initial State:**

Current location on map, time 0

**Successor Function:**

Location to every next intersection toward destination, current total time traveled to reach it.

Malik Johnson, Ph.D.

A location on the map, and the current time

Current location on map, time 0

Location to every next intersection toward destination,  
current total time traveled to reach it.

---

Matt Johnson, Ph.D.

16

---

---

---

---

---

---

## Routing Problem (2)

**Goal Test:**  
Desired address on map, shortest possible travel time.

**Path Cost:**  
Total travel time to each intermediate state

Mail Johnson, Ph.D.

17

Desired address on map, shortest possible travel time.

Total travel time to each intermediate state

Matt Johnson, Ph.D.

17

---

---

---

---

---

---

# Uniformed Search

Matt Johnson, Ph.D.

18

---

---

---

---

---

## Uninformed Search

In **uninformed search**, the algorithm has no additional information about states beyond that of the problem definition.

- Algorithm operates “blind”
- Can only generate successor states, and determine whether goal has been reached

Matt Johnson, Ph.D.

19

---

---

---

---

---

---

---

## Uninformed Search Algorithms

- Breadth-First Search
- Uniform-Cost Search
- Depth-First Search
- Depth-Limited Search
- Iterative-Deepening Search
- Bidirectional Search

Matt Johnson, Ph.D.

20

---

---

---

---

---

---

---

## Measuring Performance

### Completeness:

Is the algorithm guaranteed to find a solution?

### Optimality:

Does the strategy find the optimal solution?

### Time Complexity:

How long does it take to find a solution?

### Space Complexity:

How much memory is needed to perform the search?

Matt Johnson, Ph.D.

21

---

---

---

---

---

---

---

## Complexity Metrics

The **branching factor**  $b$  is the maximum number of successors any node can have.

The **depth**  $d$  is the depth of the shallowest goal node.

The **maximum length**  $m$  is the maximum length of any path in the state space.

Matt Johnson, Ph.D.

22

---

---

---

---

---

---

---

## Terminology

### Expanding a node (or state)

Applying the successor function to a node.

### Known State

A state that has been generated by successor function

Matt Johnson, Ph.D.

23

---

---

---

---

---

---

---

## Implementation

All search algorithms need to maintain two lists

### Closed list

- All states that have been expanded
- Needed to avoid duplication

### Frontier (or fringe or open list)

- All known states that have not yet been expanded
- The ordering of the frontier varies depending on search algorithm

Matt Johnson, Ph.D.

24

---

---

---

---

---

---

---



## Breadth-First Search (BFS)

The frontier is maintained as a queue.  
*BFS is the search for the cautious.*

Completeness: Yes

Optimality: Yes

Time Complexity:  $O(b^{d+1})$

Space Complexity:  $O(b^{d+1})$

Matt Johnson, Ph.D.

25

---

---

---


---

---

---

---

## Breadth-First Search (2)



Matt Johnson, Ph.D.

26

---

---

---

---

---

---

---

## Uniform Cost Search

Frontier is sorted by *total* path cost to reach state.  
When all path costs are 1, this is BFS

Completeness: Yes

Optimality: Yes

Time Complexity:  $O(b^{C^*/\epsilon})$

Space Complexity:  $O(b^{C^*/\epsilon})$

$C^*$  is cost of optimal solution,  $\epsilon$  is minimum action cost.

Matt Johnson, Ph.D.

27

---

---

---

---

---

---

---

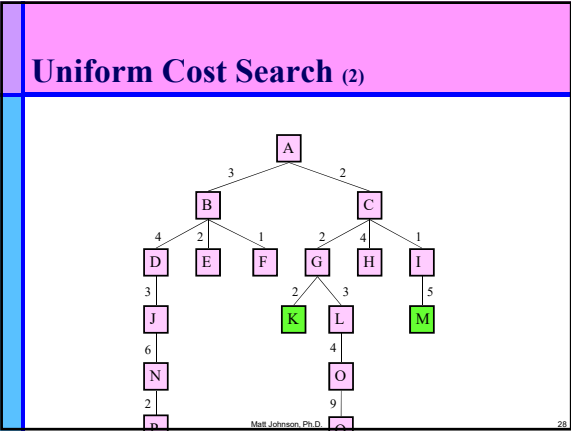
## Uniform Cost Search (2)

```

graph TD
    A[A] -- 3 --> B[B]
    A -- 2 --> C[C]
    B -- 4 --> D[D]
    B -- 2 --> E[E]
    C -- 2 --> G[G]
    C -- 4 --> H[H]
    C -- 1 --> I[I]
    D -- 3 --> J[J]
    J -- 6 --> N[N]
    N -- 2 --> P[P]
    G -- 2 --> K[K]
    G -- 3 --> L[L]
    L -- 4 --> O[O]
    O -- 9 --> Q[Q]
    H -- 5 --> M[M]
    style A fill:#f9d5e5
    style B fill:#f9d5e5
    style C fill:#f9d5e5
    style D fill:#f9d5e5
    style E fill:#f9d5e5
    style G fill:#f9d5e5
    style H fill:#f9d5e5
    style I fill:#f9d5e5
    style J fill:#f9d5e5
    style N fill:#f9d5e5
    style P fill:#f9d5e5
    style L fill:#f9d5e5
    style O fill:#f9d5e5
    style Q fill:#f9d5e5
    style K fill:#d5f9d5
    style M fill:#d5f9d5
    style Q fill:#d5f9d5
  
```

Mail Johnson, Ph.D.

28

[illegible]

## Depth-First Search (BFS)

The frontier is maintained as a stack.  
*DFS is the search for the brave.*

Completeness: No  
Optimality: No  
Time Complexity:  $O(b^m)$   
Space Complexity:  $O(bm)$

Malik Johnson, Ph.D.

29

## Depth-First Search (BFS)

The frontier is maintained as a stack.  
*DFS is the search for the brave.*

Completeness: No  
Optimality: No  
Time Complexity:  $O(b^m)$   
Space Complexity:  $O(bm)$

Malik Johnson, Ph.D.

29

## Depth-First Search (BFS)

The frontier is maintained as a stack.  
*DFS is the search for the brave.*

Completeness: No  
Optimality: No  
Time Complexity:  $O(b^m)$   
Space Complexity:  $O(bm)$

Malik Johnson, Ph.D.

29

## Depth-First Search (BFS)

The frontier is maintained as a stack.  
*DFS is the search for the brave.*

Completeness: No  
Optimality: No  
Time Complexity:  $O(b^m)$   
Space Complexity:  $O(bm)$

Malik Johnson, Ph.D.

29

## Depth-First Search (BFS)

The frontier is maintained as a stack.  
*DFS is the search for the brave.*

Completeness: No  
Optimality: No  
Time Complexity:  $O(b^m)$   
Space Complexity:  $O(bm)$

Malik Johnson, Ph.D.

29

## Depth-First Search (BFS)

The frontier is maintained as a stack.  
*DFS is the search for the brave.*

Completeness: No  
Optimality: No  
Time Complexity:  $O(b^m)$   
Space Complexity:  $O(bm)$

Malik Johnson, Ph.D.

29

---

---

---

---

---

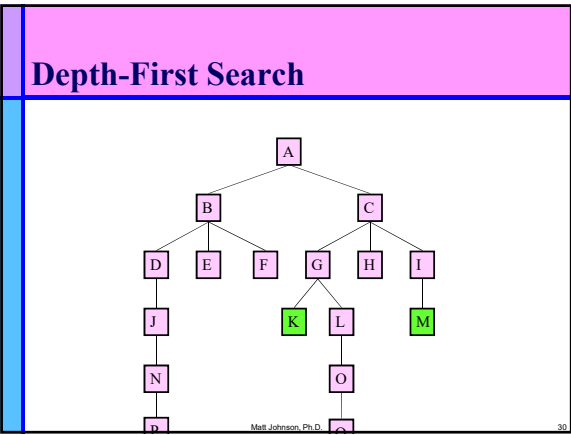
---

# Depth-First Search

```
graph TD; A[A] --> B[B]; A --> C[C]; B --> D[D]; B --> E[E]; B --> F[F]; C --> G[G]; C --> H[H]; C --> I[I]; D --> J[J]; J --> N[N]; N --> P[P]; G --> K[K]; G --> L[L]; L --> O[O]; I --> M[M]; style A fill:#fff,stroke:#f00; style B fill:#fff,stroke:#f00; style C fill:#fff,stroke:#f00; style D fill:#fff,stroke:#f00; style E fill:#fff,stroke:#f00; style F fill:#fff,stroke:#f00; style G fill:#fff,stroke:#f00; style H fill:#fff,stroke:#f00; style I fill:#fff,stroke:#f00; style J fill:#fff,stroke:#f00; style N fill:#fff,stroke:#f00; style P fill:#fff,stroke:#f00; style K fill:#0f0,stroke:#000; style L fill:#0f0,stroke:#000; style M fill:#0f0,stroke:#000; style O fill:#fff,stroke:#f00;
```

Matt Johnson, Ph.D.

30



---

---

---

---

---

---

## Depth-Limited Search

DFS with a predetermined **cutoff depth**  $l$ .  
Terminates once all nodes at  $l$  are generated.

Completeness: No

Optimality: No

Time Complexity:  $O(b^l)$

Space Complexity:  $O(bl)$

Matt Johnson, Ph.D.

31

---

---

---


---

---

---

---

## Depth-Limited Search (2)



Matt Johnson, Ph.D.

32

---

---

---

---

---

---

---

## Iterative Deepening Search

Also known as Iterative Deepening DFS.  
Depth-Limited Search with  $l = 0, 1, 2, \dots$

Completeness: Yes

Optimality: Yes

Time Complexity:  $O(b^d)$

Space Complexity:  $O(bd)$

Matt Johnson, Ph.D.

33

---

---

---


---

---

---

---

## Iterative Deepening Search (2)



Matt Johnson, Ph.D.

34

## Bidirectional Search

Two searches: one from start and one from goal.  
Often BFS both ways, or BFS/DFS combo.

Completeness: Yes\*

Optimality: Yes \*

Time Complexity:  $O(b^{d/2})$  \*

Space Complexity:  $O(b^{d/2})$  \*

\*for BFS both ways.

Matt Johnson, Ph.D.

35

## Informed Search

Matt Johnson, Ph.D.

36

## Informed Search

**Informed Search** is a strategy that uses problem-specific knowledge beyond the problem definition itself to find a solution.

A key component of informed search algorithms is a **heuristic function**.

Matt Johnson, Ph.D.

37

---

---

---

---

---

---

---

## Heuristics

A **heuristic function  $h(x)$**  (also called simply a **heuristic**) is the estimated cost of the cheapest path from node  $n$  to any goal node.

If  $x$  is a goal node,  $h(x) = 0$ .

The heuristic is used to rank alternatives in a search algorithm in order to decide which branch to follow.

Matt Johnson, Ph.D.

38

---

---

---

---

---

---

---

## Best-First Search

In **Best-First Search**, the frontier is maintained as a priority queue.

- Traditionally, the node with the lowest priority is chosen for expansion.
- Implemented like other state space searches.

Best-First Search employs an **evaluation function  $f(x)$**  to determine a node's priority.

Matt Johnson, Ph.D.

39

---

---

---

---

---

---

---

## Heuristic Search

The two main methods for Heuristic Search are:

1. Greedy Best-First Search
2. A\* Search

The difference between these two methods is the format of the evaluation function.

Matt Johnson, Ph.D.

40

---

---

---

---

---

---

---

## Greedy Best-First Search

In **Greedy Best-First Search**, the evaluation function is:

$$f(x) = h(x)$$

Only a heuristic function is used to prioritize nodes on the frontier.

Matt Johnson, Ph.D.

41

---

---

---

---

---

---

---

## Greedy Best-First Search (2)

As an example, consider the route finding problem. One straightforward heuristic to use would be

$$h(x) = \text{straight-line distance of city to goal}$$

The heuristic is useful because it has a strong correlation to actual road distance to travel.

Matt Johnson, Ph.D.

42

---

---

---

---

---

---

---

### Greedy Best-First Search (3)

Greedy Best-First Search resembles depth-first search:

Completeness: No

Optimality: No

Time Complexity:  $O(b^m)$

Space Complexity:  $O(bm)$

Matt Johnson, Ph.D.

43

---

---

---

---

---

---

---

### A\* Search

In A\* Search, the evaluation function is:

$$f(x) = g(x) + h(x)$$

- The function  $g(x)$  is the total cost from the start node to node  $x$
- The function  $h(x)$  is the estimated total cost from node  $x$  to a goal node.
- $f(x)$  is therefore the estimated cost of the cheapest solution through  $x$ .

Matt Johnson, Ph.D.

44

---

---

---


---

---

---

---

### A\* Search (2)



Matt Johnson, Ph.D.

45

---

---

---

---

---


---

---

### A\* Search (3)

For the 8-puzzle, two possible heuristics would be:

1. The number of tiles out of place.
2. The Manhattan distance of all tiles. The **Manhattan distance** of a tile is the number of actions necessary to move the tile to its correct goal position.



Matt Johnson, Ph.D.

46

---

---

---

---

---

---

---

---

### A\* Search (4)

A heuristic is **admissible** if it never overestimates the total cost to reach the goal.

A heuristic is **consistent** if, for every node  $x$  and every successor  $x'$  of  $x$ ,

$$h(x) \leq c(x) + h(x')$$

where  $c(x)$  is actual cost of reaching  $x'$  from  $x$ .

Matt Johnson, Ph.D.

47

---

---

---

---

---

---

---

---

### A\* Search (5)

**Pruning** is the process whereby a node is not expanded because its  $f(x)$  value is higher than other choices at higher depth.

Pruning can make A\* Search more efficient than uninformed search methods.

Matt Johnson, Ph.D.

48

---

---

---

---

---

---

---

---



## A\* Search (6)

Completeness: Yes, if the heuristic is admissible.

Optimality: Yes, if the heuristic is consistent.

Time Complexity:  $O(b^{d+1})$  without pruning

Space Complexity:  $O(b^{d+1})$  without pruning

Matt Johnson, Ph.D.

49

---

---

---

---

---

---

---

## Games

Matt Johnson, Ph.D.

50

---

---

---

---

---

---

---

## Why Games?

Games typically have:

- Small well-defined set of rules
- Well-defined set of knowledge
- Easy to evaluate performance
- Large state spaces
- Too large for exhaustive search methods

Matt Johnson, Ph.D.

51

---

---

---

---

---

---

---

## Game Theory

Mathematical Game Theory is a branch of Economics.

Example: **Prisoners' Dilemma**

John Nash, early 1950's

The Nash equilibrium is a decision-making theorem within game theory that states a player can achieve the desired outcome by not deviating from their initial strategy.

Matt Johnson, Ph.D.

52

---

---

---

---

---

---

---

## Game Theory (2)

Example: **Prisoners' Dilemma**

Two members of a criminal gang are arrested. The prosecutors offer each prisoner a bargain.

- If A and B each betray the other, each of them serves two years in prison
- If A betrays B but B remains silent, A will be set free and B will serve three years in prison (and vice versa)
- If A and B both remain silent, both of them will serve only one year in prison (on the lesser charge)

Matt Johnson, Ph.D.

53

---

---

---

---

---

---

---

## AI Games

In typical AI games

- There are 2 players
- Actions alternate between players
- Utility functions are equal but opposite
  - A good outcome from Player A's perspective is a bad outcome from Player B's
  - opposition between the agents' utility functions makes the situation *adversarial*

Matt Johnson, Ph.D.

54

---

---

---

---

---

---

---

## Games vs. Search

Unlike standard state space search, in adversarial search:

- Actions taken by opponents are unpredictable
- The solution to the search is a game strategy
- There may be time limits
- It's unlikely to find a direct path to goal

Matt Johnson, Ph.D.

55

---

---

---

---

---

---

---

## Adversarial Search

Matt Johnson, Ph.D.

56

---

---

---

---

---

---

---

## Adversarial Search

**Adversarial Search** is a state space search algorithm for competitive environments

- Each potential board or game position is a state
- Each possible action is a move to another state
- The state space can be HUGE!!!!!!
- Large branching factor (about 35 for chess)
- Terminal state could be deep (about 50 for chess)

Matt Johnson, Ph.D.

57

---

---

---

---

---

---

---

## Adversarial Search (2)

### Initial State:

Initial board or game configuration

### Successor Function:

All legal next actions and the resultant states

### Terminal Test:

Whether the game has entered a **terminal state** and game is won or lost

### Utility Function:

The numeric value for game outcomes  
e.g. win (1), lose (-1), draw (0)

Matt Johnson, Ph.D.

58

---

---

---

---

---

---

---

## Game Trees

The **game tree** is the search space

- Each complete path represents one possible game
- Not all paths are examined

The levels of the tree indicate a player's possible moves:

- The root of the tree is the initial state
- Next level is all of AI's moves
- Next level is all of opponent's moves

Matt Johnson, Ph.D.

59

---

---

---

---

---

---

---

## Plies

A **ply** is one turn of a game:

The **ply depth** is how many plies down the game tree are examined when choosing an action.

Example: Tic-Tac-Toe

- Root has 9 blank squares
- Ply 1: has 8 blank squares
- Ply 2: has 7 blank squares
- Ply 3: has 6 blank squares

Matt Johnson, Ph.D.

60

---

---

---


---

---

---

---

## Tic-Tac-Toe



Matt Johnson, Ph.D.

61

## Minimax Search

**Minimax Search** is a competitive search algorithm for two players agents called **MAX** and **MIN**:

**MAX**

Higher utility is better  
Starts game

**MIN**

Lower Utility is better

Matt Johnson, Ph.D.

62

## Minimax Search (2)


```
function MINIMAX-DECISION(state, game) returns an action
    action, state ← the a, s in SUCCESSORS(state)
        such that MINIMAX-VALUE(s, game) is maximized
    return action

function MINIMAX-VALUE(state, game) returns a utility value
    if TERMINAL-TEST(state) then
        return UTILITY(state)
    else if MAX is to move in state then
        return the highest MINIMAX-VALUE of SUCCESSORS(state)
    else
        return the lowest MINIMAX-VALUE of SUCCESSORS(state)
```

Matt Johnson, Ph.D.

63

## Minimax Search Example



## Ply Limit

It may not be computationally feasible to search entire game tree to all win/lose/draw.

You can instead use a **ply limit** in your search, then heuristically evaluate nodes at the limit. For Tic-Tac-Toe example, :

- $h(x) = \infty$  for win
- $h(x) = -\infty$  for loss
- $h(x) = \# \text{ of ways left to win} - \# \text{ of ways left to lose}$  otherwise

## Metrics

Complete: Yes if the tree is finite

Optimal: Yes, against an optimal opponent

Time Complexity:  $O(b^m)$

Space Complexity  $O(bm)$

## Alpha-Beta Pruning

Matt Johnson, Ph.D.

67

---

---

---

---

---

---

---

## Pruning

**Pruning** is the process whereby a node is not expanded because its  $f(x)$  value is higher/lower than other choices at higher depth.

- The problem with minimax search is that the number of game states it has to examine is exponential in the number of moves.
- Pruning can be used to reduce the time complexity of Minimax Search!

Matt Johnson, Ph.D.

68

---

---

---

---

---

---

---

## Alpha-Beta Pruning

**Alpha-Beta Pruning** remembers two values:

- $\alpha$  is the value of the best choice we have found so far for MAX at any choice point along the path
- $\beta$  is the value of the best choice we have found so far for MIN at any choice point along the path

Minimax Search can be modified to:

- prune values lower than  $\alpha$  when maximizing
- prune values higher than  $\beta$  when minimizing

Matt Johnson, Ph.D.

69

---

---

---

---

---

---

---

## Alpha-Beta Pruning (2)

```

function ALPHA-BETA-SEARCH(state, game) returns an action
  action, state  $\leftarrow$  the a, s in SUCCESSORS[game](state)
  such that MIN-VALUE(s, game,  $-\infty$ ,  $+\infty$ ) is maximized
  return action

function MAX-VALUE(state, game,  $\alpha$ ,  $\beta$ ) returns the minimax value of state
  if CUTOFF-TEST(state) then return EVAL(state)
  for each s in SUCCESSORS(state) do
     $\alpha \leftarrow \max(\alpha, \text{MIN-VALUE}(s, \text{game}, \alpha, \beta))$ 
    if  $\alpha \geq \beta$  then return  $\beta$ 
  return  $\alpha$ 

function MIN-VALUE(state, game,  $\alpha$ ,  $\beta$ ) returns the minimax value of state
  if CUTOFF-TEST(state) then return EVAL(state)
  for each s in SUCCESSORS(state) do
     $\beta \leftarrow \min(\beta, \text{MAX-VALUE}(s, \text{game}, \alpha, \beta))$ 
    if  $\beta \leq \alpha$  then return  $\alpha$ 
  return  $\beta$ 

```

Matt Johnson, Ph.D.

70

---

---

---

---

---


---

---

---

## Alpha-Beta Pruning Example

Step 1:



Matt Johnson, Ph.D.

71

---

---

---

---

---


---

---

---

## Alpha-Beta Pruning Example (2)

Step 2:



Matt Johnson, Ph.D.

72

---

---

---

---

---

---


---

---



### Alpha-Beta Pruning Example (3)

Step 3:



Matt Johnson, Ph.D.

73

---

---

---

---

---


---

---

---

### Alpha-Beta Pruning Example (4)

Step 4:



Matt Johnson, Ph.D.

74

---

---

---

---

---


---

---

---

### Alpha-Beta Pruning Example (5)

Step 5:



Matt Johnson, Ph.D.

75

---

---

---

---


---

---

---

---

## Alpha-Beta Pruning Example 2



Matt Johnson, Ph.D.

76

## Alpha-Beta Pruning Results

- Pruning *does not* affect outcome.
- Effectiveness depends on order of successors.
- If best successor is evaluated first, the search is  $O(b^{d/2})$  instead of  $O(b^d)$ .
  - This means that in the same amount of time, Alpha-Beta search can search twice as deep!
  - Can easily reach depth of 8 and play good chess (with a branching factor of ~6 instead of 35)

Matt Johnson, Ph.D.

77

## Games of Chance

Matt Johnson, Ph.D.

78

## Nondeterministic Games

A **Game of Chance** is a game with a nondeterministic environment.

- Chance is introduced by dice, card shuffling, etc.
- Probabilities of outcomes must be indicated in the game tree.

A **chance node** is a state in the game tree which is assigned a probability.

Matt Johnson, Ph.D.

79

---

---

---

---

---

---

---

## Expected Value

The **expected value** of an outcome is:

$$\sum_{i=0}^k x_i p_i$$

where  $x_i$  is the utility/heuristic value of each outcome and  $p_i$  is its chance of occurring.

Matt Johnson, Ph.D.

80

---

---

---

---

---

---

---

## Expectiminimax

**Expectiminimax** is a minimax search that handles chance nodes:

- if state is a MAX node then  
return highest Expectiminimax
- if state is a MIN node then  
return lowest Expectiminimax
- if state is a CHANCE node then  
return weighted average of Expectiminimax

Matt Johnson, Ph.D.

81

---

---

---


---

---

---

---

## Expectiminimax (2)



Mark Johnson, Ph.D.

82